

Sameer_Nepal_DSC_630_Project

November 1, 2022

```
[119]: # importing the libraries
```

```
import pandas as pd
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
```

```
[2]: #reading the dataset
```

```
df = pd.read_csv('birth.csv')
```

```
[3]: df.head(10)
```

```
[3]:   fage  mage    mature  weeks  premie  visits  gained  weight  \
0  34.0   34  younger mom    37  full term    14.0   28.0    6.96
1  36.0   31  younger mom    41  full term    12.0   41.0    8.86
2  37.0   36   mature mom    37  full term    10.0   28.0    7.51
3   NaN   16  younger mom    38  full term     NaN   29.0    6.19
4  32.0   31  younger mom    36    premie    12.0   48.0    6.75
5  32.0   26  younger mom    39  full term    14.0   45.0    6.69
6  37.0   36   mature mom    36    premie    10.0   20.0    6.13
7  29.0   24  younger mom    40  full term    13.0   65.0    6.74
8  30.0   32  younger mom    39  full term    15.0   25.0    8.94
9  29.0   26  younger mom    39  full term    11.0   22.0    9.12
```

```
   lowbirthweight  sex  habit  marital  whitemom
0         not low  male  nonsmoker  married    white
1         not low  female  nonsmoker  married    white
2         not low  female  nonsmoker  married  not white
3         not low  male  nonsmoker  not married    white
4         not low  female  nonsmoker  married    white
5         not low  female  nonsmoker  married    white
6         not low  female  nonsmoker  married    white
7         not low  male  nonsmoker  not married    white
8         not low  female  nonsmoker  married    white
```

9 not low male nonsmoker not married not white

We will look into the dataset to see which are object to convert to the numeric values

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fage                   886 non-null   float64
1   mage                   1000 non-null  int64
2   mature                 1000 non-null  object
3   weeks                  1000 non-null  int64
4   premie                 1000 non-null  object
5   visits                 944 non-null   float64
6   gained                 958 non-null   float64
7   weight                 1000 non-null  float64
8   lowbirthweight        1000 non-null  object
9   sex                    1000 non-null  object
10  habit                  981 non-null   object
11  marital                1000 non-null  object
12  whitemom               1000 non-null  object
dtypes: float64(4), int64(2), object(7)
memory usage: 101.7+ KB
```

```
[5]: from sklearn import preprocessing
from collections import defaultdict
dic = defaultdict(preprocessing.LabelEncoder)
# Encoding the categorical variable
fit = df.select_dtypes(include=['object']).apply(lambda x: dic[x.name].
    fit_transform(x))

#Converting the categorical columns using the encoding
for x in list(dic.keys()):
    df[x] = dic[x].transform(df[x])
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fage                   886 non-null   float64
1   mage                   1000 non-null  int64
2   mature                 1000 non-null  int64
3   weeks                  1000 non-null  int64
```

```

4  premie          1000 non-null   int64
5  visits          944 non-null   float64
6  gained          958 non-null   float64
7  weight          1000 non-null   float64
8  lowbirthweight  1000 non-null   int64
9  sex             1000 non-null   int64
10 habit           1000 non-null   int64
11 marital         1000 non-null   int64
12 whitemom        1000 non-null   int64
dtypes: float64(4), int64(9)
memory usage: 101.7 KB

```

```

[7]: #checking the data
df.head()

```

```

[7]:   fage  mage  mature  weeks  premie  visits  gained  weight  lowbirthweight  \
0  34.0   34      1     37      0    14.0    28.0    6.96             1
1  36.0   31      1     41      0    12.0    41.0    8.86             1
2  37.0   36      0     37      0    10.0    28.0    7.51             1
3   NaN   16      1     38      0     NaN    29.0    6.19             1
4  32.0   31      1     36      1    12.0    48.0    6.75             1

   sex  habit  marital  whitemom
0    1     0        0          1
1    0     0        0          1
2    0     0        0          0
3    1     0        1          1
4    0     0        0          1

```

```

[8]: # filling the missing data
df.fillna(df.median(), inplace = True)

```

```

[9]: df.isnull().sum()

```

```

[9]: fage          0
mage          0
mature        0
weeks         0
premie        0
visits        0
gained        0
weight        0
lowbirthweight 0
sex           0
habit         0
marital       0
whitemom      0
dtype: int64

```


We are trying to predict if the baby will be low birth or not so lowbirthweight will be our variable that we want to predict. We will create new data frame with all the variables except the dependent variable.

```
[10]: target = df['lowbirthweight']  
independent = df.drop(columns = ['lowbirthweight'])
```

Here, the target is the dataframe with just the target variable lowbirthweight and independent is the dataframe without the target variable.


```
[61]: # Distribution plot to see the skewness of the data  
import matplotlib.pyplot as plt  
fig, ax = plt.subplots(figsize =(10,6))  
sns.distplot(df.fage)
```

```
[61]: <AxesSubplot:xlabel='fage', ylabel='Density'>
```




```
[62]: fig, ax = plt.subplots(figsize =(10,6))  
sns.distplot(df.visits)
```

```
[62]: <AxesSubplot:xlabel='visits', ylabel='Density'>
```



```
[63]: fig, ax = plt.subplots(figsize=(10,6))
      sns.distplot(df.gained)
```

```
[63]: <AxesSubplot:xlabel='gained', ylabel='Density'>
```



Page, visits and gained data is little skewed so we will use median to fill the missing values.

```
[17]: # Lets look at the dataframe for features i.e independent
independent.head()
```

```
[17]:
```

	fage	mage	mature	weeks	premie	visits	gained	weight	sex	habit	\
0	34.0	34	1	37	0	14.0	28.0	6.96	1	0	
1	36.0	31	1	41	0	12.0	41.0	8.86	0	0	
2	37.0	36	0	37	0	10.0	28.0	7.51	0	0	
3	31.0	16	1	38	0	12.0	29.0	6.19	1	0	
4	32.0	31	1	36	1	12.0	48.0	6.75	0	0	

	marital	whitemom
0	0	1
1	0	1
2	0	0
3	1	1
4	0	1

We will use some techniques for feature selection

```
[19]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(independent,target)
```

```
[19]: RandomForestClassifier()
```

```
[48]: feature_imp = pd.DataFrame(rfc.feature_importances_, columns = ['RF_Value'],
↳ index = independent.columns)
feature_imp = feature_imp.reset_index()
```

```
[49]: feature_imp.sort_values(['RF_Value'], ascending = 0)
```

```
[49]:
```


	index	RF_Value
7	weight	0.675066
3	weeks	0.118978
4	premie	0.066919
5	visits	0.030589
6	gained	0.028004
1	mage	0.027362
0	fage	0.023058
9	habit	0.011839
10	marital	0.005578
8	sex	0.005341
11	whitemom	0.004059
2	mature	0.003209

Here we have different features with the importance value. Weight has the highest value because it is the weight of the baby at birth. It will be omitted from the dataset as it is showing same observation as lowbirthweight.

Looking at the correlation matrix

```
[29]: fig, ax = plt.subplots(figsize=(12,12))
      dataplot = sns.heatmap(df.corr(), cmap="YlGnBu", annot=True)

      # displaying heatmap
      plt.show()
```



Weight has high correlation and will be removed from the data because it might cause the overfitting

in the modelling.

Using the recursive feature elimination

```
[64]: from sklearn.feature_selection import RFE
      from sklearn.linear_model import LogisticRegression

      lgr = LogisticRegression()
      rfe = RFE(lgr, 8)
      fit = rfe.fit(independent, target)
```

```
[51]: selected_feature = pd.DataFrame(rfe.support_, columns = ['RFE_Selected'], index_
      => independent.columns)
      selected_feature = selected_feature.reset_index()
      selected_feature[selected_feature['RFE_Selected'] == True]
```

```
[51]:
```

	index	RFE_Selected
2	mature	True
3	weeks	True
4	premie	True
5	visits	True
7	weight	True
8	sex	True
9	habit	True
10	marital	True

Using ExtraTreesClassifier to fit a number of randomized decision trees to the data and is a from of ensemble learning

```
[52]: from sklearn.ensemble import ExtraTreesClassifier

      model1 = ExtraTreesClassifier()
      model1.fit(independent, target)
      ext = pd.DataFrame(model1.feature_importances_, columns = ["ExtraTree_Value"],
      => index=independent.columns)
      ext = ext.reset_index()
      ext.sort_values(['ExtraTree_Value'], ascending = 0)
```

```
[52]:
```

	index	ExtraTree_Value
7	weight	0.434552
4	premie	0.143632
3	weeks	0.105350
5	visits	0.059456
6	gained	0.055640
0	fage	0.055568
1	mage	0.052225
9	habit	0.030102
8	sex	0.019932
10	marital	0.018870

11	whitemom	0.012909
2	mature	0.011764

Using Lasso Regression for feature selection

```
[53]: from sklearn.svm import LinearSVC
from sklearn.feature_selection import SelectFromModel
lsvc = LinearSVC(C=0.01, penalty="l1", dual=False).fit(independent, target)
model2 = SelectFromModel(lsvc, prefit=True)
l1 = pd.DataFrame(model2.get_support(), columns = ["L1"], index=independent.
↳ columns)
l1 = l1.reset_index()
l1[l1['L1'] == True]
```

```
[53]:      index  L1
0     fage  True
1     mage  True
3    weeks  True
6   gained  True
7   weight  True
```

Combining the results together

```
[55]: from functools import reduce
all_df = [feature_imp, selected_feature, ext, l1]
result = reduce(lambda left, right: pd.merge(left, right, on='index'), all_df)
result
```

```
[55]:      index  RF_Value  RFE_Selected  ExtraTree_Value  L1
0     fage  0.023058         False         0.055568  True
1     mage  0.027362         False         0.052225  True
2    mature  0.003209          True         0.011764 False
3    weeks  0.118978          True         0.105350  True
4    premie  0.066919          True         0.143632 False
5    visits  0.030589          True         0.059456 False
6    gained  0.028004         False         0.055640  True
7    weight  0.675066          True         0.434552  True
8       sex  0.005341          True         0.019932 False
9    habit  0.011839          True         0.030102 False
10  marital  0.005578          True         0.018870 False
11 whitemom  0.004059         False         0.012909 False
```

Constructing a table for total score

```
[65]: columns = ['RF_Value', 'ExtraTree_Value']
table = pd.DataFrame({}, [])
table['index'] = result['index']
for x in columns:
```

```

    table[x] = result['index'].isin(list(result.nlargest(6,x)['index'])).
    ↳astype(int)

table['RFE_Selected'] = result['RFE_Selected'].astype(int)
table['L1'] = result['L1'].astype(int)

table['Total'] = table.sum(axis=1)
table.sort_values('Total', ascending =0)

```

```

[65]:
      index  RF_Value  ExtraTree_Value  RFE_Selected  L1  Total
3    weeks         1             1             1    1     4
7   weight         1             1             1    1     4
4   premie         1             1             1    0     3
5  visits         1             1             1    0     3
6   gained         1             1             0    1     3
0     fage         0             1             0    1     2
1     mage         1             0             0    1     2
2   mature         0             0             1    0     1
8      sex         0             0             1    0     1
9    habit         0             0             1    0     1
10  marital         0             0             1    0     1
11 whitemom        0             0             0    0     0

```

The above table shows what the important feature are that are related to the lowbirthweight. My aim was to find if the smoking habit cause the low bith weight but after looking at the table construted above we can tell that habit have low contribution to the low birth weight and other facotors are involved in this process.

I will be now looking into creating the model that predicts if the baby will be low bith weight on the basis of different factors present in the data set but will be dropping whitemom and marital column form the dataset.

```

[66]: #Looking at our dataframe
df.head()

```

```

[66]:
   fage  mage  mature  weeks  premie  visits  gained  weight  lowbirthweight  \
0  34.0   34      1     37      0     14.0    28.0    6.96             1
1  36.0   31      1     41      0     12.0    41.0    8.86             1
2  37.0   36      0     37      0     10.0    28.0    7.51             1
3  31.0   16      1     38      0     12.0    29.0    6.19             1
4  32.0   31      1     36      1     12.0    48.0    6.75             1

   sex  habit  marital  whitemom
0    1     0        0          1
1    0     0        0          1
2    0     0        0          0

```

```

3      1      0      1      1
4      0      0      0      1

```

```
[69]: #dropping column marital and whitemom after performing feature selection
df.drop(["marital","whitemom"], axis =1, inplace = True)
```

```
[70]: df
```

```
[70]:
```

	fage	mage	mature	weeks	premie	visits	gained	weight \
0	34.0	34	1	37	0	14.0	28.0	6.96
1	36.0	31	1	41	0	12.0	41.0	8.86
2	37.0	36	0	37	0	10.0	28.0	7.51
3	31.0	16	1	38	0	12.0	29.0	6.19
4	32.0	31	1	36	1	12.0	48.0	6.75
..
995	28.0	24	1	39	0	12.0	20.0	6.49
996	37.0	31	1	38	0	8.0	33.0	5.80
997	27.0	27	1	34	1	7.0	25.0	6.75
998	31.0	33	1	42	0	0.0	13.0	7.44
999	21.0	17	1	41	0	12.0	41.0	7.13

	lowbirthweight	sex	habit
0	1	1	0
1	1	0	0
2	1	0	0
3	1	1	0
4	1	0	0
..
995	1	1	2
996	1	0	0
997	1	0	0
998	1	1	0
999	1	0	0

```
[1000 rows x 11 columns]
```

- Now we will use the remaining features to develop our model

```
[81]: # splitting the data in training and test
X = df.drop(columns = ['lowbirthweight'])
y = df['lowbirthweight']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=60,
↳train_size = .8)
print( '---Shape of Splits---')
print('X train:{}'.format(X_train.shape))
print('X test:{}'.format(X_test.shape))
print('y train:{}'.format(y_train.shape))
```

```
print('y test:{}'.format(y_test.shape))
```

---Shape of Splits---

X train:(800, 10)

X test:(200, 10)

y train:(800,)

y test:(200,)

Scaling the data before feeding to the model. It helps to reduce the effect of the outlier in the model's prediction.

```
[127]: from sklearn.preprocessing import StandardScaler
```

```
# scaling the input data
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)
```

Developing the model using the final dataset for test and train

1. Logistic Regression Model

```
[88]: # Building the Logistic Regression Model
model_1 = LogisticRegression()
model_1.fit(X_train,y_train)
y_pred = model_1.predict(X_test)

# testing the model and displaying the score
score = model_1.score(X_test,y_test)
print(score)
```

0.99

```
[90]: # looking at the accuracy score of the model
accuracy = accuracy_score(y_test, y_pred)
print('Accuracy score: ', accuracy)
```

Accuracy score: 0.99


```
[139]: # calculating the Confusion Matrix
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# Plot the confusion matrix in graph
cm = confusion_matrix(y_test,y_pred)

# plotting with labels
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot()

# showing the matrix
```

```
plt.show()
```



Classification report is used to measure the quality of predictions from a classification algorithm. We will display the classification report for the model.

```
[107]: # displaying the classification report  
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.87	0.93	15
1	0.99	1.00	0.99	185
accuracy			0.99	200
macro avg	0.99	0.93	0.96	200
weighted avg	0.99	0.99	0.99	200

2. RandomForestClassifier Model

```
[116]: # creating a RF classifier  
clf = RandomForestClassifier(n_estimators = 100)  
  
# fit function is used to train the model using the training sets as parameters  
clf.fit(X_train, y_train)
```

```

# performing predictions on the test dataset
y_predic = clf.predict(X_test)

# metrics are used to find accuracy or error
from sklearn import metrics
print()

# using metrics module for accuracy calculation
print("ACCURACY OF THE MODEL: ", metrics.accuracy_score(y_test, y_predic))

```

ACCURACY OF THE MODEL: 0.995

```


[131]: import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# Plot the confusion matrix in graph
cm = confusion_matrix(y_test, y_predic)

# plotting with labels
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot()

# showing the matrix
plt.show()

```



```
[132]: # Looking at the model classification report
print(classification_report(y_test, y_predic))
```

	precision	recall	f1-score	support
0	1.00	0.93	0.97	15
1	0.99	1.00	1.00	185
accuracy			0.99	200
macro avg	1.00	0.97	0.98	200
weighted avg	1.00	0.99	0.99	200

3. KNN Classification Model

```
[141]: # creating the model
from sklearn.neighbors import KNeighborsClassifier


knn = KNeighborsClassifier(n_neighbors=1)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
ac = metrics.accuracy_score(y_test, y_pred)
print('Accuracy score of model is', ac)

# plotting confusion matrix with labels
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot()

# showing the matrix
plt.show()

# Looking at the model classification report
print(classification_report(y_test, y_pred))
```

Accuracy score of model is 0.97



	precision	recall	f1-score	support
0	0.91	0.67	0.77	15
1	0.97	0.99	0.98	185
accuracy			0.97	200
macro avg	0.94	0.83	0.88	200
weighted avg	0.97	0.97	0.97	200

[]: